

Arijit Gupta

(+44) 7342 887938 | arijitgupta2000@gmail.com | [linkedin.com/in/arijitgupta42](https://www.linkedin.com/in/arijitgupta42)

PROFILE

AI Engineer with nearly 3 years of experience in LLM orchestration, agentic systems, retrieval-augmented generation, and production deployment of generative-AI applications, working embedded with enterprise customers from problem definition to delivery. Seeking full-time roles in Forward Deployed Engineering, Applied AI, and related areas.

WORK EXPERIENCE

Cognizant

London, UK

AI Engineer

Feb 2026 – Present

- Sole AI Engineer on an agentic data-discovery tool on Azure, letting analysts self-serve queries across 200+ enterprise datasets in plain English and cutting discovery from days of back-and-forth with the statistics team to minutes.
- Designed the orchestration layer (Azure OpenAI, FastAPI/ASGI): a coordinating LLM decomposes each query, routes sub-tasks to specialised LLM calls or Python functions, then consolidates the results – shipped as an Azure Function with response streaming, graceful failure handling, and managed-identity auth, integrated into the client's existing infrastructure.
- Embedded with analyst end users, translating ambiguous business-language requests into concrete technical requirements through rapid demo-and-feedback cycles.

ZS Associates

Bangalore, India

Advanced Data Science Associate

Jul 2022 – Aug 2024

- Built and productionized an LLM chatbot (LangChain, Haystack) answering insurance agents' plain-English questions – from single-policy lookups to geographic aggregates – by compiling each into a dynamically generated Python retrieval pipeline, increasing the client calls each agent could handle by 25%.
- Engineered the data-acquisition backend powering the chatbot: a scraping-and-parsing pipeline that structured 3,000+ unstructured insurance policies into a queryable knowledge base.
- Architected an LLM-driven insight-mining pipeline over 100K+ unstructured healthcare documents, combining clustering with RAG-based entity extraction across 10K+ physician notes to cut retrieval latency 40% and inform safety reviews and marketing strategy.

Microsoft

Hyderabad, India

Software Engineer Intern

May 2021 – Jul 2021

- Developed voice assistant features for Microsoft Search (TypeScript, Express), improving enterprise search accessibility and daily user engagement.
- Built RESTful APIs, frontend components, and YAML conversational datasets with product and design teams, improving intent recognition and accelerating delivery of production features.

TECHNICAL SKILLS

Languages & Tools: Python, SQL, TypeScript, FastAPI, Git

LLM & GenAI: Agentic Orchestration, RAG, LLM Evaluation, Azure OpenAI, LangChain, Hugging Face, PyTorch

Cloud & Deployment: Azure (Functions, Container Apps, AI Search, Managed Identity/RBAC), AWS (Lambda, SageMaker, IAM)

EDUCATION

The University of Sheffield

Sheffield, UK

MSc Computer Science with Speech and Language Processing | Distinction

Sep 2024 – Nov 2025

Dissertation: Hallucination Span Detection in Large Language Models – benchmarked open-source LLMs against proprietary models for hallucination span detection; found them competitive on detection and highly effective as context retrievers, with end-to-end pipeline performance dropping only 1–2% when proprietary retrieval was replaced.

BITS Pilani

Goa, India

B.E. Electronics and Instrumentation

Aug 2018 – May 2022

PUBLICATIONS & HONORS

Towards Using Diachronic Distributed Word Representations as Models of Lexical Development – Gupta, Patil, Baths.

ACL 2022 Student Research Workshop (poster).

First Runner-Up, Web/Mobile category, Global PyTorch Summer Hackathon 2020 (2,500+ participants) – *Rasoe*: self-scraped 308-class food dataset, EfficientNet-B2 TorchScript-optimised for on-device inference, shipped as web and Android apps. Featured in PyTorch's official publication.